

lightgbm-GDBT

July 13, 2019

```
In [11]: import numpy as np
import pandas as pd
from sklearn import preprocessing
import gc, sys
gc.enable()

import lightgbm as lgb

In [6]: def feature_engineering(is_train=True, debug=True):
    test_idx = None
    if is_train:
        print("processing train.csv")
        if debug == True:
            df = pd.read_csv('./train_V2.csv', nrows=10000)
        else:
            df = pd.read_csv('./train_V2.csv')

        df = df[df['maxPlace'] > 1]
    else:
        print("processing test.csv")
        df = pd.read_csv('./test_V2.csv')
        test_idx = df.Id

    print("remove some columns")
    target = 'winPlacePerc'

    print("Adding Features")

    df['headshotrate'] = df['kills']/df['headshotKills']
    df['killStreakrate'] = df['killStreaks']/df['kills']
    df['healthitems'] = df['heals'] + df['boosts']
    df['totalDistance'] = df['rideDistance'] + df["walkDistance"] + df["swimDistance"]
    df['killPlace_over_maxPlace'] = df['killPlace'] / df['maxPlace']
    df['headshotKills_over_kills'] = df['headshotKills'] / df['kills']
    df['distance_over_weapons'] = df['totalDistance'] / df['weaponsAcquired']
    df['walkDistance_over_heals'] = df['walkDistance'] / df['heals']
    df['walkDistance_over_kills'] = df['walkDistance'] / df['kills']
```

```

df['killsPerWalkDistance'] = df['kills'] / df['walkDistance']
df["skill"] = df["headshotKills"] + df["roadKills"]

df[df == np.Inf] = np.NaN
df[df == np.NINF] = np.NaN

print("Removing Na's From DF")
df.fillna(0, inplace=True)

features = list(df.columns)
features.remove("Id")
features.remove("matchId")
features.remove("groupId")
features.remove("matchType")

# matchType = pd.get_dummies(df['matchType'])
# df = df.join(matchType)

y = None

if is_train:
    print("get target")
    y = np.array(df.groupby(['matchId', 'groupId'])[target].agg('mean'), dtype=np.float64)
    features.remove(target)

print("get group mean feature")
agg = df.groupby(['matchId', 'groupId'])[features].agg('mean')
agg_rank = agg.groupby('matchId')[features].rank(pct=True).reset_index()

if is_train: df_out = agg.reset_index()[['matchId', 'groupId']]
else: df_out = df[['matchId', 'groupId']]

df_out = df_out.merge(agg.reset_index(), suffixes=["", ""], how='left', on=['matchId', 'groupId'])
df_out = df_out.merge(agg_rank, suffixes=["_mean", "_mean_rank"], how='left', on=['matchId', 'groupId'])

print("get group max feature")
agg = df.groupby(['matchId', 'groupId'])[features].agg('max')
agg_rank = agg.groupby('matchId')[features].rank(pct=True).reset_index()
df_out = df_out.merge(agg.reset_index(), suffixes=["", ""], how='left', on=['matchId', 'groupId'])
df_out = df_out.merge(agg_rank, suffixes=["_max", "_max_rank"], how='left', on=['matchId', 'groupId'])

print("get group min feature")
agg = df.groupby(['matchId', 'groupId'])[features].agg('min')
agg_rank = agg.groupby('matchId')[features].rank(pct=True).reset_index()
df_out = df_out.merge(agg.reset_index(), suffixes=["", ""], how='left', on=['matchId', 'groupId'])
df_out = df_out.merge(agg_rank, suffixes=["_min", "_min_rank"], how='left', on=['matchId', 'groupId'])

```

```

print("get group size feature")
agg = df.groupby(['matchId', 'groupId']).size().reset_index(name='group_size')
df_out = df_out.merge(agg, how='left', on=['matchId', 'groupId'])

print("get match mean feature")
agg = df.groupby(['matchId'])[features].agg('mean').reset_index()
df_out = df_out.merge(agg, suffixes=["", "_match_mean"], how='left', on=['matchId'])

print("get match size feature")
agg = df.groupby(['matchId']).size().reset_index(name='match_size')
df_out = df_out.merge(agg, how='left', on=['matchId'])

df_out.drop(["matchId", "groupId"], axis=1, inplace=True)

X = df_out

feature_names = list(df_out.columns)

del df, df_out, agg, agg_rank
gc.collect()

return X, y, feature_names, test_idx

```

```
In [5]: x_train, y_train, train_columns, _ = feature_engineering(True, False)
```

```

processing train.csv
remove some columns
Adding Features
Removing Na's From DF
get target
get group mean feature
get group max feature
get group min feature
get group size feature
get match mean feature
get match size feature

```

```
In [7]: x_test, _, _ , test_idx = feature_engineering(False, True)
```

```

processing test.csv
remove some columns
Adding Features
Removing Na's From DF
get group mean feature
get group max feature
get group min feature
get group size feature

```

```
get match mean feature
get match size feature
```

```
In [8]: # Thanks and credited to https://www.kaggle.com/gemartin who created this wonderful mem
def reduce_mem_usage(df):
    """ iterate through all the columns of a dataframe and modify the data type
        to reduce memory usage.
    """
    start_mem = df.memory_usage().sum()
    print('Memory usage of dataframe is {:.2f} MB'.format(start_mem))

    for col in df.columns:
        col_type = df[col].dtype

        if col_type != object:
            c_min = df[col].min()
            c_max = df[col].max()
            if str(col_type)[:3] == 'int':
                if c_min > np.iinfo(np.int8).min and c_max < np.iinfo(np.int8).max:
                    df[col] = df[col].astype(np.int8)
                elif c_min > np.iinfo(np.int16).min and c_max < np.iinfo(np.int16).max:
                    df[col] = df[col].astype(np.int16)
                elif c_min > np.iinfo(np.int32).min and c_max < np.iinfo(np.int32).max:
                    df[col] = df[col].astype(np.int32)
                elif c_min > np.iinfo(np.int64).min and c_max < np.iinfo(np.int64).max:
                    df[col] = df[col].astype(np.int64)
            else:
                if c_min > np.finfo(np.float16).min and c_max < np.finfo(np.float16).max:
                    df[col] = df[col].astype(np.float16)
                elif c_min > np.finfo(np.float32).min and c_max < np.finfo(np.float32).max:
                    df[col] = df[col].astype(np.float32)
                else:
                    df[col] = df[col].astype(np.float64)
        else:
            df[col] = df[col].astype('category')

    end_mem = df.memory_usage().sum()
    print('Memory usage after optimization is: {:.2f} MB'.format(end_mem))
    print('Decreased by {:.1f}%'.format(100 * (start_mem - end_mem) / start_mem))

    return df

x_train = reduce_mem_usage(x_train)
x_test = reduce_mem_usage(x_test)
```

```
Memory usage of dataframe is 4021060096.00 MB
Memory usage after optimization is: 948516192.00 MB
```

Decreased by 76.4%
Memory usage of dataframe is 3837401216.00 MB
Memory usage after optimization is: 903259258.00 MB
Decreased by 76.5%

```
In [12]: #excluded_features = []
        #use_cols = [col for col in df_train.columns if col not in excluded_features]

        train_index = round(int(x_train.shape[0]*0.8))
        dev_X = x_train[:train_index]
        val_X = x_train[train_index:]
        dev_y = y_train[:train_index]
        val_y = y_train[train_index:]
        gc.collect();

        # custom function to run light gbm model
        def run_lgb(train_X, train_y, val_X, val_y, x_test):
            params = {"objective" : "regression", "metric" : "mae", 'n_estimators':20000, 'early_stopping_rounds':100,
                      "num_leaves" : 31, "learning_rate" : 0.05, "bagging_fraction" : 0.7,
                      "bagging_seed" : 0, "num_threads" : 4, "colsample_bytree" : 0.7}

            lgtrain = lgb.Dataset(train_X, label=train_y)
            lgval = lgb.Dataset(val_X, label=val_y)
            model = lgb.train(params, lgtrain, valid_sets=[lgtrain, lgval], early_stopping_rounds=100)

            pred_test_y = model.predict(x_test, num_iteration=model.best_iteration)
            return pred_test_y, model

        # Training the model #
        pred_test, model = run_lgb(dev_X, dev_y, val_X, val_y, x_test)

/home/sirius/anaconda3/lib/python3.6/site-packages/lightgbm/engine.py:118: UserWarning: Found `num_leaves` in params. Will use it instead of argument.
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))
/home/sirius/anaconda3/lib/python3.6/site-packages/lightgbm/engine.py:123: UserWarning: Found `num_threads` in params. Will use it instead of argument.
warnings.warn("Found `{}` in params. Will use it instead of argument".format(alias))

Training until validation scores don't improve for 200 rounds.
[1000]    training's l1: 0.0282507    valid_1's l1: 0.0288693
[2000]    training's l1: 0.0270028    valid_1's l1: 0.0281773
[3000]    training's l1: 0.026192     valid_1's l1: 0.0278832
[4000]    training's l1: 0.0255316    valid_1's l1: 0.0277071
[5000]    training's l1: 0.0249701    valid_1's l1: 0.0275958
[6000]    training's l1: 0.0244602    valid_1's l1: 0.0275161
[7000]    training's l1: 0.0240017    valid_1's l1: 0.0274589
[8000]    training's l1: 0.0235657    valid_1's l1: 0.0274077
```

[9000]	training's l1: 0.0231617	valid_1's l1: 0.0273648
[10000]	training's l1: 0.0227831	valid_1's l1: 0.0273305
[11000]	training's l1: 0.022418	valid_1's l1: 0.0273009
[12000]	training's l1: 0.0220675	valid_1's l1: 0.0272781
[13000]	training's l1: 0.0217335	valid_1's l1: 0.0272615
[14000]	training's l1: 0.021404	valid_1's l1: 0.0272413
[15000]	training's l1: 0.0210929	valid_1's l1: 0.0272243
[16000]	training's l1: 0.0207937	valid_1's l1: 0.0272085
[17000]	training's l1: 0.0205064	valid_1's l1: 0.0271988
[18000]	training's l1: 0.0202239	valid_1's l1: 0.0271857
[19000]	training's l1: 0.019947	valid_1's l1: 0.0271731
[20000]	training's l1: 0.0196855	valid_1's l1: 0.0271666

Did not meet early stopping. Best iteration is:

[20000]	training's l1: 0.0196855	valid_1's l1: 0.0271666
---------	--------------------------	-------------------------